

Mustakim Shikalgar

☎ 480-328-4635 ✉ shikalgar.mustakim@gmail.com 🌐 mustakim-shikalgar.dev 📄 linkedin.com/in/mustakim-shikalgar

EXPERIENCE

Full Stack Developer, GenAI Team

Jun 2026 – Present

Saayam for All (open-source nonprofit)

Remote

- Reverse-engineered the as-built architecture of a 7-team serverless platform (7 API Gateway APIs, 209 routes, 200 Lambda functions) from live infrastructure exports after finding existing documentation described a system that was never built; verified every figure independently in the AWS console
- Identified 33 deployed API routes with no authorizer, including DELETE and PUT operations on user data addressable by path parameter alone, with no resource policy mitigating the exposure
- Uncovered two parallel implementations of the same service serving live traffic under one namespace, one undeployed for 12 months and missing a year of fixes, by cross-referencing gateway integration targets against the deployment pipeline
- Found provider API credentials stored in plaintext Lambda environment variables outside the organization's credential rotation and unread by any code path; scoped their removal and rotation

Software Engineer (Industry-Sponsored Capstone)

Jan – May 2026, Sep 2026 – Present

METY Legal AI Platform, MyEdMaster

Tempe, AZ

- Owned end-to-end backend architecture for a three-service legal AI platform (React, Django, FastAPI/LangGraph) on a 6-person team across 7 Agile sprints; Django owns all persistence while a stateless AI service owns every LLM call with no database access
- Cut per-query LLM cost 82% (\$0.024 → \$0.0044) by removing an unconditional gpt-4o call found through LangSmith profiling, folding clarification, analysis, topic detection, and confidence scoring into one structured call, and routing classification to gpt-4o-mini
- Built end-to-end request tracing by propagating an X-Request-ID from Django through every AI-service log line, with structured JSON logs that record counts, scores, and durations but never prompts or user documents
- Built async conversation summarization as a Celery task on Redis (3 retries) that triggers past 1,500 tokens and 6 messages and runs after the response returns, so users never wait on it

SELECTED SOFTWARE ENGINEERING PROJECTS

AegisFlow | Python, FastAPI, Kubernetes, Redis, pgvector, NATS JetStream, OpenTelemetry

Feb – May 2026

- Architected a 7-microservice reliability platform (14 containers, 9,500+ lines) with DAG-based orchestration and a 5-component confidence model routing outputs to accept, repair, fallback, or reject
- Built atomic Redis-based rate limiting, circuit breakers, an event-sourced replay engine, and hybrid retrieval; validated reliability via a dedicated chaos engineering service, instrumented with 11 Prometheus metric families and OpenTelemetry

Distributed Key-Value Store | Java 17, Raft Consensus, gRPC, Protocol Buffers, Docker

Jan 2024 – Dec 2025

- Implemented Raft consensus from scratch across a 5-node cluster: leader election (150-300ms randomized timeouts), log replication, and quorum-based commit, with tunable per-request consistency (linearizable vs. eventual reads)
- Benchmarked 17,857 ops/sec write throughput and sub-1ms p99 read latency, validating zero data loss under simultaneous 2-of-5 node failure via automated fault-injection testing

TECHNICAL SKILLS

Languages: Java, Python, TypeScript, JavaScript, SQL

Systems & Networking: gRPC, Protocol Buffers, REST APIs, asyncio

Cloud & DevOps: AWS (Lambda, SSM Parameter Store, S3), Docker, Kubernetes, Git, CI/CD, Prometheus, OpenTelemetry

Data: PostgreSQL, MongoDB, Redis, pgvector, Qdrant

Backend Frameworks: FastAPI, Django, Celery, LangGraph, LangChain

EDUCATION

Arizona State University

Tempe, AZ

Master of Science in Software Engineering, Data Science Specialization, GPA: 3.77/4.0

May 2026

ACHIEVEMENTS

LeetCode: 900+ problems solved | Knight badge | Contest rating 2000+ | Top 3% globally | 140+ Hard

Publication: IEEE COMPSAC 2025, Missing Persons Knowledge Graph (27% acceptance rate)